

ASHRAE - Great Energy Predictor III

Hyperparameter Optimization and Stacking

Siddharth
Yadav

Agenda

kaggle
DAYS

MEETUP

- About me
- Kaggle Kernels - Why you should use them?
- About the ASHRAE competition
- Current trends and tips for the competition
- Using Bayesian Optimization - Why? When? How?
- Stacking - The rewards of blending
- Q&A

About me

kaggle
DAYS

MEETUP

- Elitist nerd due to circumstances
- B.Tech Final Year Student
- Kaggle Kernels Master
- Passionate for all things “Data”
- Got into Data Science just before the hype
- Worked for organisations like



INTERNSHALA

kaggle™



Devfolio

Kaggle Kernels

Jupyter Notebooks
that run on a
hosted runtime on
kaggle free of
charge.

You don't have to
worry about
storage,
processing or other
resources.

You also get GPU
and version control
on the platform
itself. importantly,

Most important
component of any
platform: good
community
support.

Why use kernels?

Github works for developers. but how to showcase your skills in analytics and statistics.

Kernels provide more exposure and feedback.

High end kernels make good projects for your resume. Helped me nail an interview. True story.

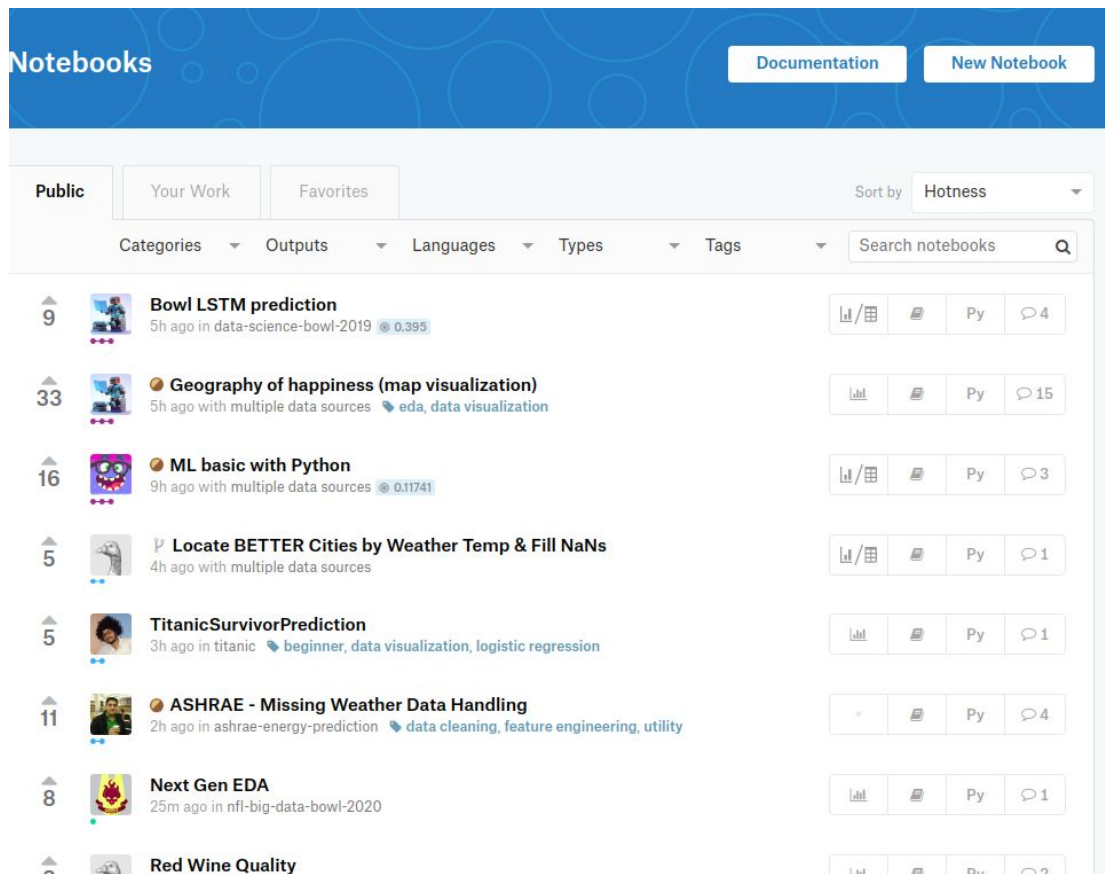
The stuff you learn through kernels won't be found in books or online courses.

How to approach kernels?

Easy. Explore them.

Analytics Competitions:

- ✓ Real data science.
- ✓ Data storytelling.
- ✓ Real Life Solutions
- ✓ Prize Money 💰



The screenshot displays the Kaggle Notebooks interface. At the top, there's a blue header with the word "Notebooks" and buttons for "Documentation" and "New Notebook". Below this, a navigation bar includes tabs for "Public", "Your Work", and "Favorites", along with a "Sort by" dropdown set to "Hotness". A search bar labeled "Search notebooks" is also present. The main content area lists several notebooks with their titles, creation times, and associated tags. Each notebook entry includes a small icon, a rank (e.g., 9, 33, 16), and a set of icons for viewing the notebook, its code, and comments.

Rank	Notebook Title	Time Ago	Tags	Views	Code	Comments
9	Bowl LSTM prediction	5h ago in data-science-bowl-2019	0.395	1	Py	4
33	Geography of happiness (map visualization)	5h ago with multiple data sources	eda, data visualization	1	Py	15
16	ML basic with Python	9h ago with multiple data sources	0.11741	1	Py	3
5	Locate BETTER Cities by Weather Temp & Fill NaNs	4h ago with multiple data sources		1	Py	1
5	TitanicSurvivorPrediction	3h ago in titanic	beginner, data visualization, logistic regression	1	Py	1
11	ASHRAE - Missing Weather Data Handling	2h ago in ashrae-energy-prediction	data cleaning, feature engineering, utility		Py	4
8	Next Gen EDA	25m ago in nfl-big-data-bowl-2020		1	Py	1
	Red Wine Quality				Py	?

INTERNSHALA

kaggle™

LOGICA

Devfolio

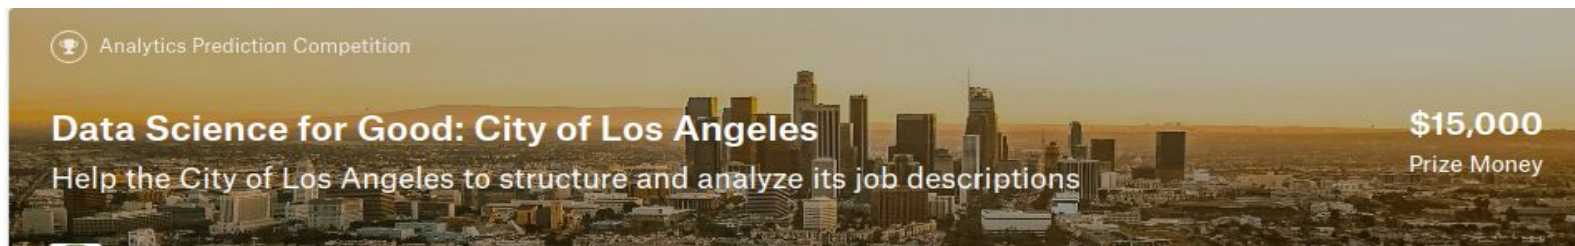



 Dataset

Data Science for Good: DonorsChoose.org

Help DonorsChoose.org connect donors with projects they care about

 563

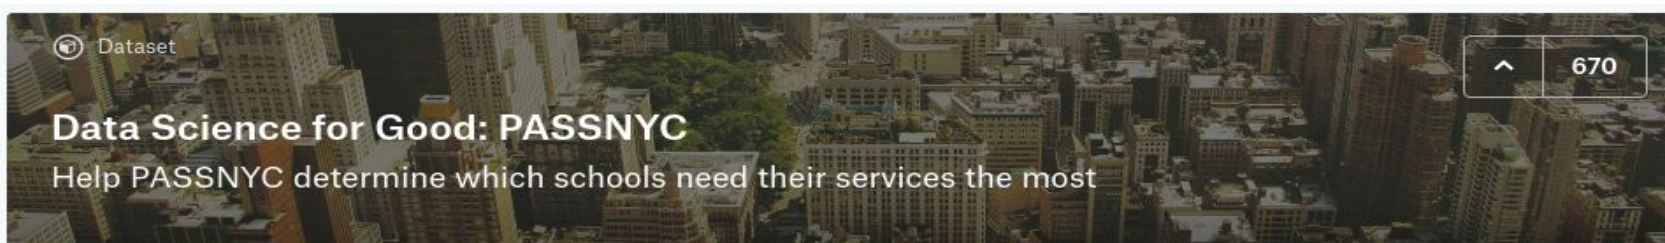



 Analytics Prediction Competition

Data Science for Good: City of Los Angeles

Help the City of Los Angeles to structure and analyze its job descriptions

\$15,000
Prize Money



 Dataset

Data Science for Good: PASSNYC

Help PASSNYC determine which schools need their services the most

670

Create kernels that
impart valuable
knowledge.

Share insights that
help the data science
community.



Why I think kaggle kernels is a semi-successful product?

Colab launched around the same time and provides almost all functionalities except version control and “Commit & Chill”.



Colab is mainstream.

Kaggle has put some restrictions on kernels

Illegal
cryptocurrency
mining.



Exploitation of
hotness algorithm
to gain more
upvotes



GPU usage is free
but not unlimited



The new ASHRAE competition

Basically, an American professional organisation of HVAC engineers wants you to predict metered building energy usage in the following areas: chilled water, electric, hot water, and steam meters with millions of rows of training data and even more test data.

THE DATA

Training Data	Over 20 million rows
Test Data	Over 41.6 million rows
Weather data for training set	Around 139,000 rows
Weather data for test set	Around 277,000 rows
Buildings metadata	Around 1450 rows

Current trends and tips for the competition

Using Bayesian Optimization

Why? When? How?

The Different Methods of Optimization

Manual

It is all about guessing and intuition.

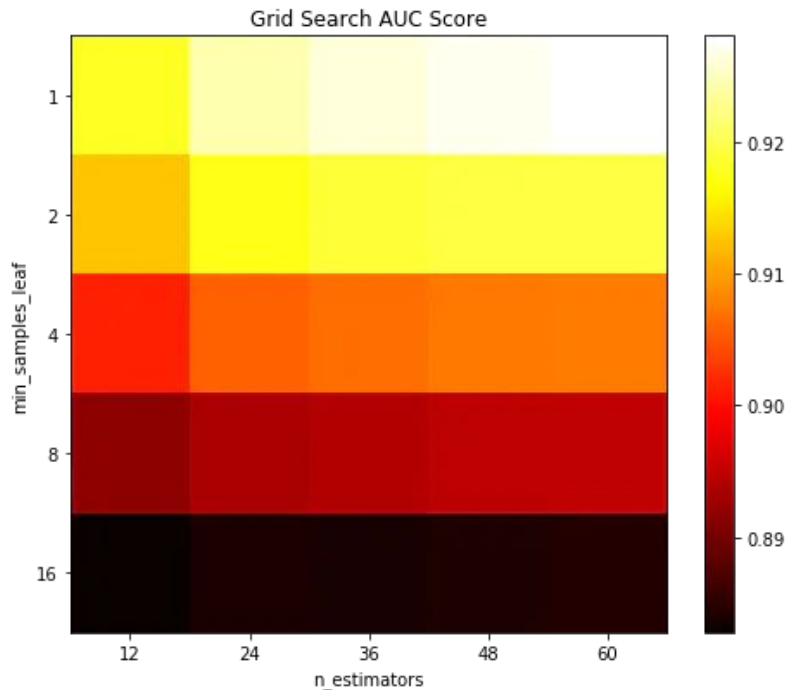
Random
Search

Model is iteratively trained on random parameters by the random search algorithms.


The different methods of optimization

Grid Search

A grid of parameters is formed. Each combination of parameters is used to train models by the grid search algorithm. The best combination of parameters according to the given metric is finally selected.



Grid Layout



Random Layout




Image from <http://jmlr.csail.mit.edu/papers/volume13/bergstra12a/bergstra12a.pdf>

```
svc = SVC()  
svc.get_params()
```

```
param_grid =  
[{'C': [0.01, 1, 10, 100, 1000],  
 'kernel': ['linear', 'poly', 'rbf', 's  
igmoid'],  
 'gamma': [0.01, 1, 10, 100, 1000]}]
```

The different methods of optimization

kaggle
DAYS

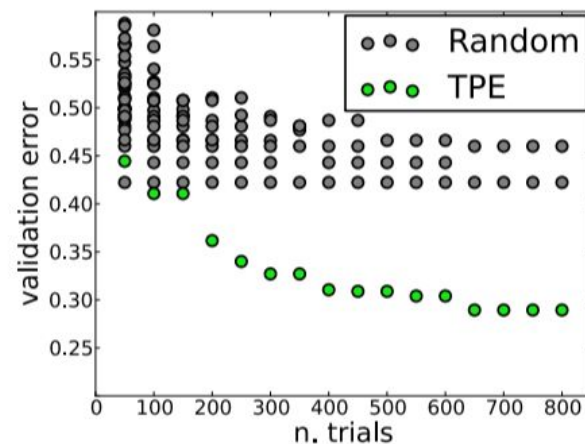

MEETUP

Bayesian Optimization

The one-sentence summary of Bayesian hyperparameter optimization is: build a probability model of the objective function and use it to select the most promising hyperparameters to evaluate in the true objective function.

Parameters of the next iteration are decided by a probabilistic model which keeps track of past evaluations.

$$P(\text{score} \mid \text{hyperparameters})$$



Steps to follow

Define an objective function

This is a function that takes your data, trains the desired model on it and compares the validation set with predictions for the given metric. The metric's value is returned.

Steps to follow

Define a search space

Usually a dictionary containing either arrays of values for a parameter or a range of values for a parameter (can be a distribution)

Steps to follow

Run trials to
get best
parameters

Run trials/iterations(at least 15-16) using the objective function for the given search space. Eventually the best trial results are given as output.

When do we use this?

kaggle
DAYS

MEETUP

Works for large datasets when grid search space is too computationally expensive causing memory errors. High number of trials may not be feasible given large data size.

Surpasses grid search in terms of speed and performance for small datasets and simpler models almost always.

Results on ASHRAE

1.20-1.26 score
on various single
models

Bayesian
Optimization
with 20 trials



1.18 on single
LightGBM model

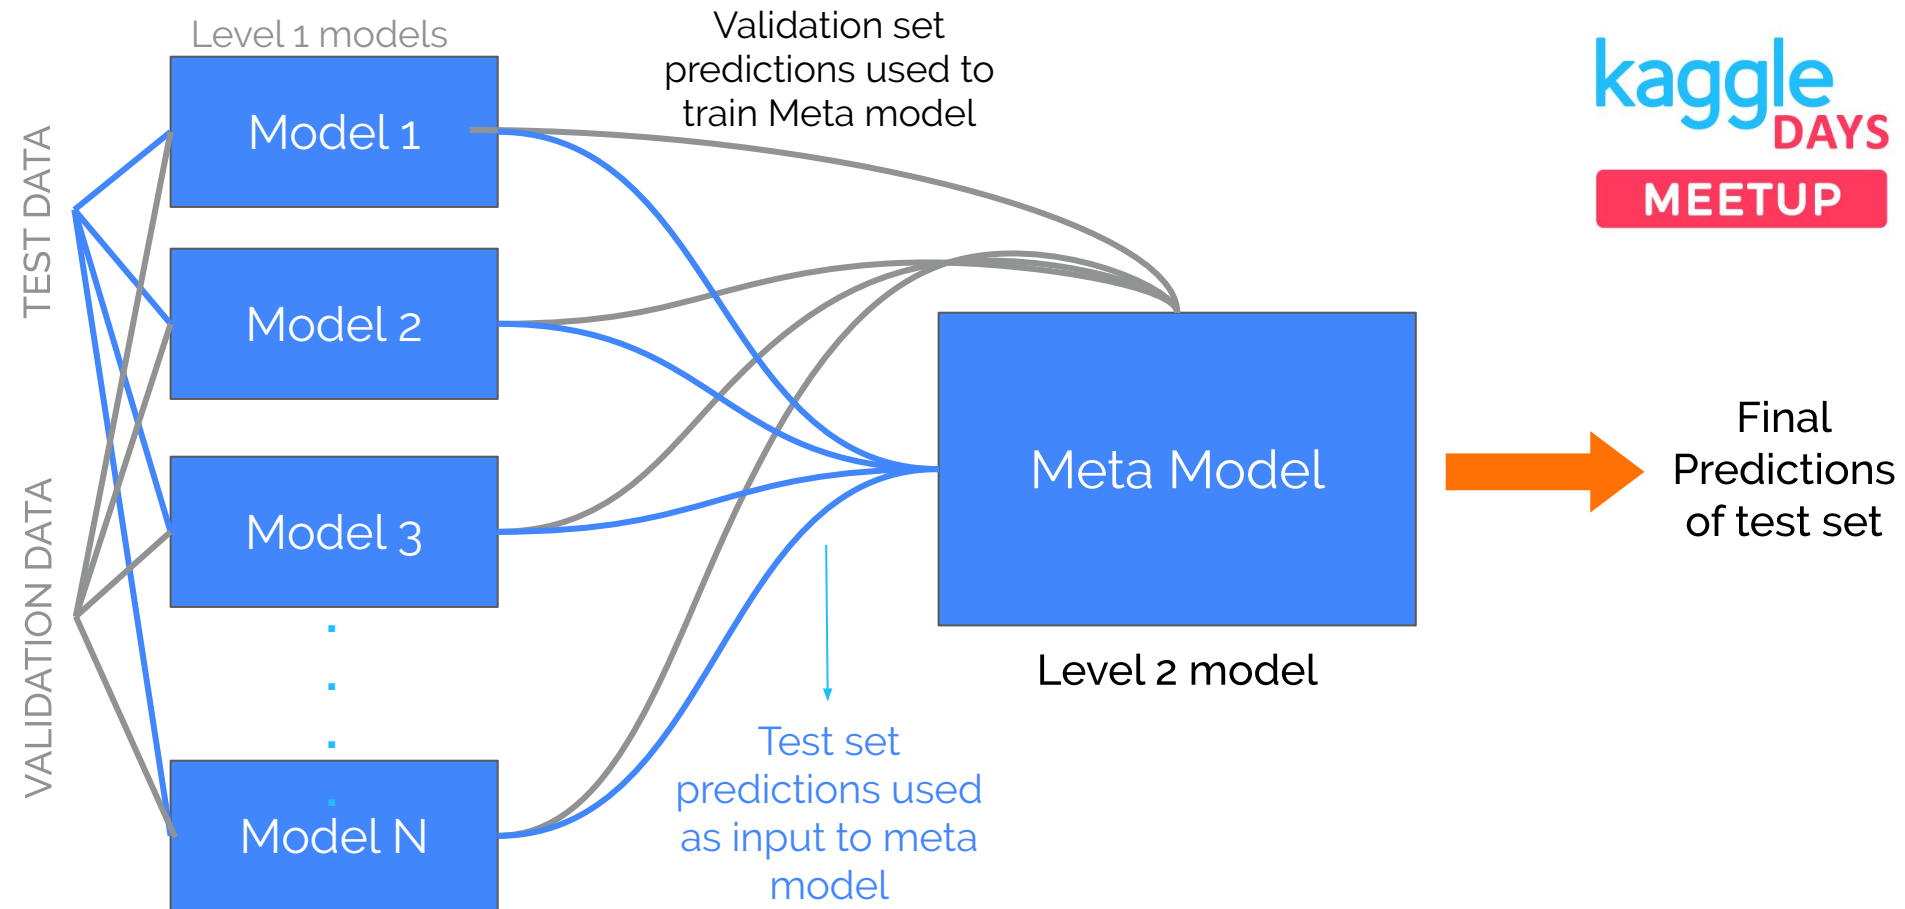
Better than single models and performance similar
to ensemble model of 15 folds(RMSLE 1.18)

Stacking

The rewards of blending

It is an ensembling method where you are using predictions of multiple models on validation set(first level) against real targets to train a new model(second level).

This second level model is then used to get final predictions using first level predictions of test set as its input data.



Results on ASHRAE

1.18 using 15
folds LGBM
ensemble

Stacking



1.14 using 10 folds of
LGBM+CatBoost with
3-layer Neural Network
as Meta Regressor

NN Model is non-deterministic so different
predictions each time.

Inspirations and References

1. <https://qiita.com/bmjo114/items/8009f282c99b77780563>
2. <https://blog.floydhub.com/guide-to-hyperparameters-search-for-deep-learning-models/>
3. <https://datafai.com/hyper-parameters-optimization-using-gridsearch-and-crossvalidation/>
4. <https://towardsdatascience.com/a-conceptual-explanation-of-bayesian-model-based-hyperparameter-optimization-for-machine-learning-b8172278050f>

Thanks for your time